

PHASE 1 — IMMEDIATE READINESS (Weeks 1–4)

🎯 Narrative Focus

"I'm a modern web architect who understands frontend at scale and is pragmatically leading the AI transition."

◆ What to KEEP (Validated)

Topic	What to Focus On (Refined)	Resources
React as architecture	Whiteboard-level understanding only (don't read everything)	https://react.dev/learn/render-and-commit
Concurrent rendering	Why it exists, when it helps/hurts	https://react.dev/reference/react/useTransition
Micro-frontends	Why they fail more often than they succeed	https://martinfowler.com/articles/micro-frontends.html
Node event loop	Diagram + explanation only (most common senior gotcha)	https://nodejs.org/en/learn/asynchronous-work/event-loop
AI executive vocabulary	Cost, latency, hallucinations, guardrails	https://www.databricks.com/glossary/retrieval-augmented-generation

◆ NEW (Added from Feedback)

Gap Filled	Why It Matters at Your Level	Resource
LLM mechanics vocabulary	Expert expect fluency, not math	https://platform.openai.com/docs/guides/text-generation
Context window / tokens	Cost & scalability conversations	https://www.anyscale.com/blog/llm-context-windows
Temperature & determinism	Risk & predictability in prod	https://docs.anthropic.com/claude/docs/temperature

◆ Project Step (Refined)

Step	Action
Project Step 1 Node.js script calling OpenAI or AWS Bedrock → “Hello AI”	
Goal	Remove fear, gain vocabulary, prove hands-on credibility

Step	Action
Timebox	1–2 evenings

PHASE 2 — SYSTEM DESIGN & CLOUD (Weeks 5–8)

🎯 Narrative Focus

"I design systems that scale, are observable, and don't bankrupt the company."

◆ What to KEEP (Validated)

Topic	Focus	Resources
Multi-region architecture	Conceptual design, not configs	https://aws.amazon.com/architecture/multi-region/
CDN + caching	Draw cache layers	https://aws.amazon.com/caching/
BFF pattern	Frontend-driven backend design	https://samnewman.io/patterns/architectural/bff/
Rate limiting	Protecting systems at scale	https://cloud.google.com/architecture/rate-limiting-strategies-techniques

◆ NEW (Critical Gap Filled)

Topic	Why Leaders Must Know This	Resource
Infrastructure as Code (IaC)	Modern cloud is unreadable without it	https://martinfowler.com/articles/infrastructure-as-code.html
Terraform / CDK (conceptual)	Team-scale infra management	https://developer.hashicorp.com/terraform/intro
Back-of-envelope math	System design	https://www.educative.io/blog/back-of-the-envelope-calculations

◆ Project Step (Optimized)

Step	Action
Project Step 2 Add Vector DB (Pinecone / OpenSearch) → enable RAG	
Cloud tie-in	Deploy via AWS (Lambda / API Gateway)
Outcome	AI + Cloud + Node in one system

PHASE 3 — LEADERSHIP LAYER (Weeks 9–12)

🎯 Narrative Focus

“I can lead teams, manage risk, and make AI decisions responsibly.”

◆ What to KEEP (Validated)

Topic	Focus	Resource
Monolith vs microservices	When NOT to split	https://martinfowler.com/articles/monolith-first.html
Schema evolution	Safe change management	https://www.confluent.io/blog/schema-evolution/
Eventual consistency	Trade-offs	https://www.allthingsdistributed.com/2008/12/eventually_consistent.html

◆ NEW (Major Leadership Upgrade)

Topic	Why This Matters	Resource
End-to-End MLOps pipeline	Leaders must visualize flow	https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning
AI risk management	Regulatory & reputational risk	https://www.nist.gov/ai
OSS vs Proprietary models	Strategic decision-making	https://www.anyscale.com/blog/open-source-vs-closed-langs

◆ Behavioral Preparation (Added)

Task	Details
STAR stories (5)	Conflict, failure, mentoring, influence, AI decision
Executive reflection	“What I’d do differently” (critical at senior level)

PHASE 4 — MASTERY & VISION (Months 3–4)

Narrative Focus

“AI adoption is 80% systems and people, 20% models.”

◆ What to KEEP (Validated)

Topic	Focus	Resource
Strangler pattern	AI in legacy systems	https://martinfowler.com/articles/strangler-application.html
Build vs buy	Long-term cost & control	https://martinfowler.com/articles/build-or-buy.html
AI observability	Trust & reliability	https://opentelemetry.io/docs/

◆ NEW (Executive-Level Additions)

Topic	Why Expert Care	Resource
AI governance	Board-level concern	https://cloud.google.com/architecture/ml-model-governance
Cost controls	CFO alignment	https://www.anyscale.com/blog/llm-cost-optimization
Regulation readiness	Future-proofing	https://www.oecd.org/ai/

FINAL REFINEMENT (Key Takeaways)

-  The original plan was **correct**, but too dense
-  This version:
 - Prevents burnout
 - Combines AI + AWS + Node into **one capstone**
 - Elevates you from “*hands-on senior*” to “*AI-capable technical leader*”
-  You should **not** aim to read every link
-  Aim to **explain, draw, and defend trade-offs**

Primary Capstone (Final Decision)

Serverless RAG Bot on AWS

- React chat UI
- Node.js Lambda backend
- AWS Bedrock / OpenAI

- Vector DB
- Observability + cost awareness

This single project **covers 80% of your surface area.**

expectations	https://staffeng.com/guides/
Engineering leadership vs IC depth	https://martinfowler.com/articles/leadership.html
Making trade-offs at scale	https://martinfowler.com/articles/decision-making.html

PHASE 1 — IMMEDIATE READINESS

◆ React — Senior / Principal Level

Topic	Learning Resources
React rendering architecture	https://react.dev/learn/render-and-commit
Concurrent rendering	https://react.dev/reference/react/useTransition
Suspense	https://react.dev/reference/react/Suspense
Streaming SSR	https://react.dev/reference/react-dom/server
Server vs Client Components	https://react.dev/learn/server-components
Web Vitals (LCP, CLS, TTI)	https://web.dev/vitals/
Frontend performance at scale	https://web.dev/fast/
Bundle splitting strategies	https://web.dev/code-splitting/
Hydration cost	https://web.dev/rendering-on-the-web/
Micro-frontends (when/why)	https://martinfowler.com/articles/micro-frontends.html
Module Federation	https://webpack.js.org/concepts/module-federation/
State management at scale	https://react.dev/learn/scaling-up-with-reducer-and-context
Domain-driven frontend state	https://martinfowler.com/articles/modularizing-react-apps.html
Frontend observability	https://opentelemetry.io/docs/concepts/observability-primer/
Error budgets	https://sre.google/sre-book/error-budgets/

◆ Node.js — Senior Systems Thinking

Topic	Learning Resources
Event loop (high level)	https://nodejs.org/en/learn/asynchronous-work/event-loop

Topic	Learning Resources
Async execution model	https://nodejs.org/en/docs/guides/event-loop-timers-and-nexttick/
Backpressure & streams	https://nodejs.org/en/learn/modulesstreams
CPU vs IO workloads	https://nodejs.org/en/learn/asynchronous-work/blocking-vs-non-blocking
Horizontal scaling	https://nodejs.org/en/learn/asynchronous-work/scaling-nodejs
API versioning	https://martinfowler.com/articles/api-versioning.html
Caching strategies	https://aws.amazon.com/caching/
Failure modes & retries	https://martinfowler.com/articles/patterns-of-distributed-systems/retry.html

◆ AI/ML — Executive-Level Literacy

Topic	Learning Resources
Supervised vs unsupervised	https://developers.google.com/machine-learning/intro
Training vs inference	https://cloud.google.com/architecture/ml-inference
Model drift	https://cloud.google.com/architecture/mlops-continuous-delivery-and-automation-pipelines-in-machine-learning
Hallucinations	https://www.anthropic.com/research
Precision / Recall	https://developers.google.com/machine-learning/crash-course/classification/precision-and-recall
RAG vs fine-tuning	https://www.databricks.com/glossary/retrieval-augmented-generation
Cost-latency–accuracy trade-offs	https://www.anyscale.com/blog
Guardrails	https://platform.openai.com/docs/guides/safety-best-practices
Human-in-the-loop	https://cloud.google.com/architecture/human-in-the-loop

◆ Low-Lift ML Project — Content Classification

Topic	Learning Resources
Text classification basics	https://scikit-learn.org/stable/tutorial/text_analytics/working_with_text_data.html
Transformers API (simple)	https://huggingface.co/docs/transformers/tasks/sequence_classification
Model lifecycle	https://ml-ops.org/content/model-lifecycle
Serving ML via API	https://fastapi.tiangolo.com/

Topic	Learning Resources
Node ↔ Python integration	https://martinfowler.com/articles/patterns-of-distributed-systems/api-gateway.html

PHASE 2 — CLOUD & SYSTEM DESIGN

◆ AWS + Cloud Architecture

Topic	Learning Resources
Multi-region architecture	https://aws.amazon.com/architecture/multi-region/
Caching layers	https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/Introduction.html
Event-driven systems	https://aws.amazon.com/event-driven-architecture/
API Gateway vs LB	https://aws.amazon.com/compare/the-difference-between-api-gateway-and-load-balancer/
Async workflows	https://aws.amazon.com/serverless/event-driven-architecture/
Cost optimization	https://aws.amazon.com/aws-cost-management/
Security boundaries	https://docs.aws.amazon.com/wellarchitected/latest/security-pillar/welcome.html

◆ Frontend + Cloud System Design

Topic	Learning Resources
Global frontend architecture	https://web.dev/distribute-content-efficiently/
Backend-for-Frontend (BFF)	https://samnewman.io/patterns/architectural/bff/
AI feature rollout & fallback	https://martinfowler.com/articles/feature-toggles.html

◆ High-ROI AI Project #1 — RAG System

Topic	Learning Resources
RAG architecture	https://www.pinecone.io/learn/retrieval-augmented-generation/
LangChain	https://python.langchain.com/docs/
Vector databases	https://www.pinecone.io/learn/vector-database/

Topic	Learning Resources
Evaluation of RAG	https://www.anyscale.com/blog/evaluating-llm-applications
Source attribution	https://docs.langchain.com/docs/use_cases/question_answering/citations

PHASE 3 — SYSTEM DESIGN DEPTH

Topic	Learning Resources
Monolith vs microservices	https://martinfowler.com/articles/monolith-first.html
SQL vs NoSQL	https://aws.amazon.com/nosql/
Data modeling	https://martinfowler.com/articles/evodb.html
Eventual consistency	https://www.allthingsdistributed.com/2008/12/eventually_consistent.html
Rate limiting	https://cloud.google.com/architecture/rate-limiting-strategies-techniques
Backward compatibility	https://martinfowler.com/articles/consumerDrivenContracts.html
Schema evolution	https://www.confluent.io/blog/schema-evolution/

PHASE 4 — MASTERY & VISION

◆ AI in Legacy Systems

Topic	Learning Resources
Incremental adoption	https://martinfowler.com/articles/strangler-fight-application.html
API-first AI	https://cloud.google.com/architecture/api-first
Risk mitigation	https://www.nist.gov/ai

◆ Leading AI Teams

Topic	Learning Resources
AI team composition	https://www.thoughtworks.com/insights/articles/building-effective-ai-teams
Model governance	https://cloud.google.com/architecture/ml-model-governance
Cost controls	https://www.anyscale.com/blog/llm-cost-optimization
Evaluation frameworks	https://www.evidentlyai.com/

◆ Long-Term AI Strategy

Topic	Learning Resources
Build vs buy	https://martinfowler.com/articles/build-or-buy.html
Data moat	https://stratechery.com/
AI observability	https://opentelemetry.io/docs/
Regulation readiness	https://www.oecd.org/ai/

CRITICAL AREAS (DO NOT SKIP)

Topic	Learning Resources
DSA (reasoning focus)	https://neetcode.io/roadmap (conceptual use only)
OOAD / UML	https://www.uml-diagrams.org/
Database design	https://use-the-index-luke.com/
Concurrency concepts	https://queue.acm.org/detail.cfm?id=2745385
Security basics	https://owasp.org/www-project-top-ten/
Cost modeling	https://aws.amazon.com/well-architected/
Trade-off articulation	https://martinfowler.com/articles/architecture-decision-records.html